

# FINAL REPORT

Music Genre Classification

## *Group 96*

So Hee Yoon

1004327604

Ho Seok (David) Lee

1004112177

Myeong Hun (David) Song

1004815961

Ting Wei (Sherman) Lin

1004835413

Word count: 2251

# 1. Introduction

Classifying music by ear is a fairly trivial human task; you can show a child one of Chopin's études and they will be able to recognize it as a classical piece. Formulating explicit musical criteria that defines a classical piece, however, is very difficult and requires detailed descriptions of musical properties such as instrumentation, tempo, and arrangement. Even with these descriptions, identification of a genre without listening to the music can become ambiguous. Despite these shortcomings, we observed that these explicit musical criteria reflect a pattern-like consistency within genres that machines can possibly learn to detect.

An image recognition algorithm can detect a dog by recognizing features that are distinctive of a canine. These machine-learned features, however, cannot easily be translatable for humans to understand. In the same manner, algorithms can learn some mechanized version of those explicit musical criteria.

An algorithm that can classify music is extremely helpful for music platforms such as Spotify or Apple Music. Users are more inclined to search for and explore through genres instead of groupings like "songs with guitars", since these are classifications that can not only be easily understood by humans, but are often reflective of one's music preference. Considering all of these factors, we believe that a machine that can recognize genres is both viable and valuable.

## 2. Illustration / Figure

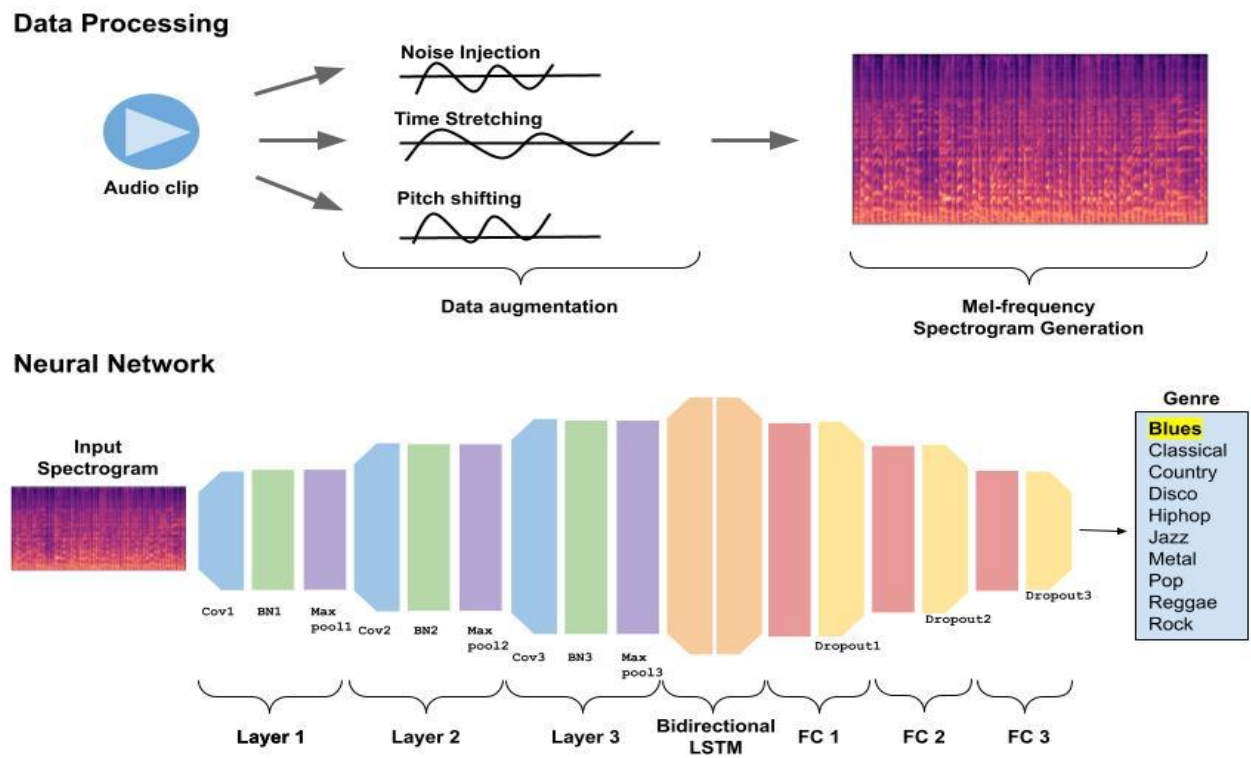


Figure 1: Illustration showing the core idea and architecture of the project

### 3. Background & Related Work

Analyzing lyrics and chords to classify music genre *by Timothy Greer and Shrikanth Narayanan [1]*

This paper uses natural language processing to classify music from the Billboard charts and represent the relationship between chord progression and lyric sequence in predicting genres.

Music genre classification *by Derek A. Huang [2]*

Convolutional Neural Networks are used to predict genres of the music, out of 10 distinct music genres. The model inputs two different data: raw amplitude data and their transformed mel-spectrograms to determine which input resulted in a higher accuracy.


## 4. Data Processing

### *Source of data*

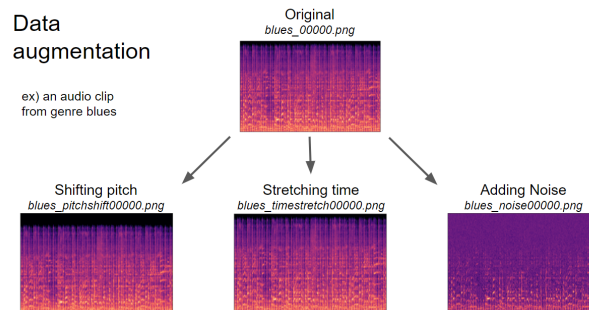
We decided to use the GTZAN genre dataset [3], which feature clips of music (of file type .wav) of 30 seconds in length of categorized genres: 10 genres with 100 audio clips in each. Data augmentation tripled the dataset by generating three new clips from each through noise injection, pitch shift, and time stretching.

### *Mel-Spectrogram*

1. Load audio file using (librosa.load)
2. Create mono audio signal (librosa.to\_mono)
3. Generate mel spectrogram (librosa.feature.melspectrogram)
4. Convert power spectrogram into decibel units (librosa.power\_to\_db)
5. Save figures as image file in the genre folder
6. Resize the 1200\*800 pixel image to 240\*160 pixel



*Figure 2: Example of generated spectrogram from audio file in blues genre.*



*Figure 3: Spectrograms generated from an audio clip through data augmentation.*

# Mel-Frequency Cepstrum Coefficient

- 1) Load audio file using (librosa.load)
- 2) Extract mfcc features (librosa.feature.mfcc)
- 3) Save mfcc values into data.csv file

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63	64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95	96	97	98	99	100	101	102	103	104	105	106	107	108	109	110	111	112	113	114	115	116	117	118	119	120	121	122	123	124	125	126	127	128	129	130	131	132	133	134	135	136	137	138	139	140	141	142	143	144	145	146	147	148	149	150	151	152	153	154	155	156	157	158	159	160	161	162	163	164	165	166	167	168	169	170	171	172	173	174	175	176	177	178	179	180	181	182	183	184	185	186	187	188	189	190	191	192	193	194	195	196	197	198	199	200	201	202	203	204	205	206	207	208	209	210	211	212	213	214	215	216	217	218	219	220	221	222	223	224	225	226	227	228	229	230	231	232	233	234	235	236	237	238	239	240	241	242	243	244	245	246	247	248	249	250	251	252	253	254	255	256	257	258	259	260	261	262	263	264	265	266	267	268	269	270	271	272	273	274	275	276	277	278	279	280	281	282	283	284	285	286	287	288	289	290	291	292	293	294	295	296	297	298	299	300	301	302	303	304	305	306	307	308	309	310	311	312	313	314	315	316	317	318	319	320	321	322	323	324	325	326	327	328	329	330	331	332	333	334	335	336	337	338	339	340	341	342	343	344	345	346	347	348	349	350	351	352	353	354	355	356	357	358	359	360	361	362	363	364	365	366	367	368	369	370	371	372	373	374	375	376	377	378	379	380	381	382	383	384	385	386	387	388	389	390	391	392	393	394	395	396	397	398	399	400	401	402	403	404	405	406	407	408	409	410	411	412	413	414	415	416	417	418	419	420	421	422	423	424	425	426	427	428	429	430	431	432	433	434	435	436	437	438	439	440	441	442	443	444	445	446	447	448	449	450	451	452	453	454	455	456	457	458	459	460	461	462	463	464	465	466	467	468	469	470	471	472	473	474	475	476	477	478	479	480	481	482	483	484	485	486	487	488	489	490	491	492	493	494	495	496	497	498	499	500	501	502	503	504	505	506	507	508	509	510	511	512	513	514	515	516	517	518	519	520	521	522	523	524	525	526	527	528	529	530	531	532	533	534	535	536	537	538	539	540	541	542	543	544	545	546	547	548	549	550	551	552	553	554	555	556	557	558	559	560	561	562	563	564	565	566	567	568	569	570	571	572	573	574	575	576	577	578	579	580	581	582	583	584	585	586	587	588	589	590	591	592	593	594	595	596	597	598	599	600	601	602	603	604	605	606	607	608	609	610	611	612	613	614	615	616	617	618	619	620	621	622	623	624	625	626	627	628	629	630	631	632	633	634	635	636	637	638	639	640	641	642	643	644	645	646	647	648	649	650	651	652	653	654	655	656	657	658	659	660	661	662	663	664	665	666	667	668	669	670	671	672	673	674	675	676	677	678	679	680	681	682	683	684	685	686	687	688	689	690	691	692	693	694	695	696	697	698	699	700	701	702	703	704	705	706	707	708	709	710	711	712	713	714	715	716	717	718	719	720	721	722	723	724	725	726	727	728	729	730	731	732	733	734	735	736	737	738	739	740	741	742	743	744	745	746	747	748	749	750	751	752	753	754	755	756	757	758	759	760	761	762	763	764	765	766	767	768	769	770	771	772	773	774	775	776	777	778	779	780	781	782	783	784	785	786	787	788	789	790	791	792	793	794	795	796	797	798	799	800	801	802	803	804	805	806	807	808	809	810	811	812	813	814	815	816	817	818	819	820	821	822	823	824	825	826	827	828	829	830	831	832	833	834	835	836	837	838	839	840	841	842	843	844	845	846	847	848	849	850	851	852	853	854	855	856	857	858	859	860	861	862	863	864	865	866	867	868	869	870	871	872	873	874	875	876	877	878	879	880	881	882	883	884	885	886	887	888	889	890	891	892	893	894	895	896	897	898	899	900	901	902	903	904	905	906	907	908	909	910	911	912	913	914	915	916	917	918	919	920	921	922	923	924	925	926	927	928	929	930	931	932	933	934	935	936	937	938	939	940	941	942	943	944	945	946	947	948	949	950	951	952	953	954	955	956	957	958	959	960	961	962	963	964	965	966	967	968	969	970	971	972	973	974	975	976	977	978	979	980	981	982	983	984	985	986	987	988	989	990	991	992	993	994	995	996	997	998	999	1000
--	---	---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	------

Figure 4: Example of generated MFCC(mel-frequency cepstrum coefficient) logged in excel sheet.

## Statistics

### Data distribution

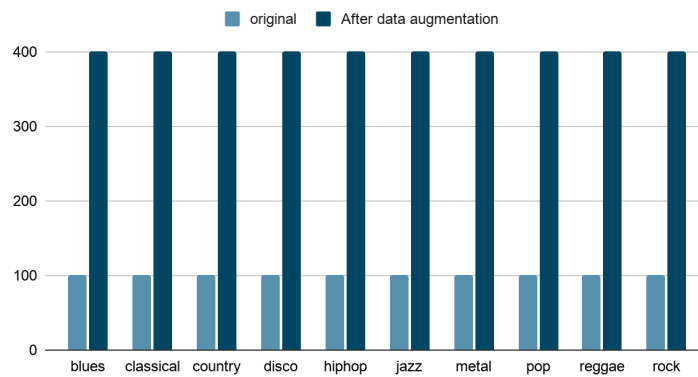


Figure 5: Cleaned data distribution for each genre.

### Total Number of dataset

- 400 audio clips in each genre
- Total of 4000 audio clips

### Dataset split ratio

- 70% training data (2800 audio clips)
- 15% validation data (600 audio clips)
- 15% test data (600 audio clips)

## 5. Architecture

Between the option of training with mel-spectrogram images or MFCC numbers, we believe that mel-spectrograms offer more possible information for the machine to learn from. Building upon this, we decided to use a convolutional neural network with a fully-connected classification neural network with ReLU activation functions.

### *Convolutional Neural Network*

There are three increasing convolutional neural networks layers channel sizes of 32, 64, and 128 respectively, each with a kernel size of 3 and stride and padding of 1. After each convolutional layer, we decided to perform batch normalization to improve training efficiency and maximum pooling layers to help consolidate learned information.

### *Fully-Connected Neural Network*

The classifier is a simple fully-connected neural network with hidden units of 128, 64, and 10 respectively, with dropout layers of 50% to avoid overfitting to the training data. The last output layer represents the classifications for each of the genres, and must be applied cross entropy loss.

### *Long Short-Term Memory*

After noticing that mel-spectrograms not only contained information about the frequencies but also the order in which they appear, or the arrangement of the song, we decided to incorporate a bidirectional long short-term memory module between the CNN and the classification ANN. The feature maps extracted through the CNN are fed in sequence to the LSTM, which are converted into the embedding for the classifier. We decided to use 128 units in the hidden layer, as making the model too complex resulted in overfitting or slower training. This allows the overall model to learn as much about the mel-spectrograms as possible.

## 6. Baseline Model

To compare the effectiveness of our model, we will compare the accuracy with a few rudimentary supervised baseline models. The two baseline models that we will be using are the  $k$ -nearest neighbours algorithm ( $k$ -NN) and a support-vector machine (SVM). Both of these models will be implemented in the scikit-learn Python library [4][5].

### *k-Nearest Neighbours Algorithm*

The  $k$ -nearest neighbours algorithm is a classification algorithm that processes labeled data [6]. Given an unlabeled data point, the algorithm finds the closest labeled data point within the parameter vector space.

Each mel-spectrogram from the labeled training data will be converted into 20 mel-frequency cepstral coefficients (MFCC) to populate the vector space. For each unobserved datum, using the Minkowski distance

$$dist(x, z) = \left( \sum_{r=1}^d |x_r - z_r|^p \right)^{1/p},$$

the algorithm will find 5 nearest neighbours (i.e. closest sounding clips). The predicted label will be the most commonly occurring label from the 5 neighbours, weighted by their distance.

### *Support Vector Machine*

A support vector machine attempts to maximize the separation between clusters of labeled vectors [7]. Since we do not know the general shape and behaviour of our data, we can blindly assume that the classes are not linearly separable. Changing the SVM kernel from a linear function to the Gaussian radial basis function (RBF)

$$rbf(x, z) = \exp(-\gamma ||x - z||^2),$$

allows the vector space to be mapped to an infinite dimensional feature space for the SVM to create non-linear decision boundaries. Once trained, the SVM can begin predicting the labels of unobserved data.



## 7. Quantitative Results

The KNN algorithm had a testing accuracy of 66.85%, and validation accuracy of 44.67%. The SVM had a testing accuracy of 47.57%, and validation accuracy of 42.67%, which are lower than K-nearest Neighbours algorithm. We have decided to use the KNN algorithm performance as the baseline model.

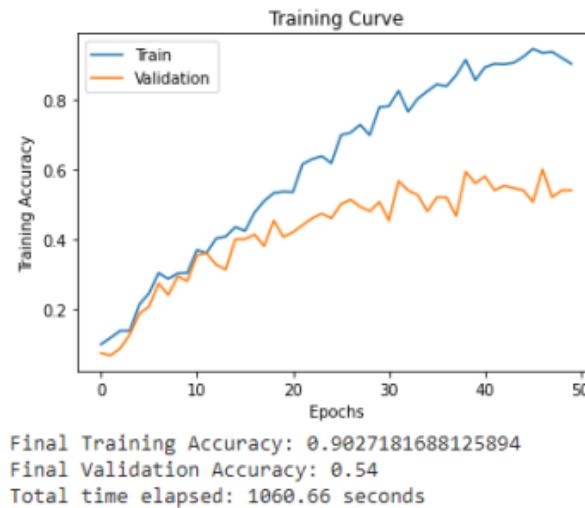


Figure 6: Training and validation accuracy curves with 1000 samples, 50 epochs, learning rate  $1 \times 10^{-4}$ .

The first trial run had 1000 samples in dataset, and the model had 50 epochs with Adam optimizer which had learning rate of  $1 \times 10^{-4}$ . We obtained training accuracy of 90%, but the validation accuracy was only 54%, indicating overfitting.





Figure 7: Training and validation accuracy curves with 4000 samples, 50 epochs, learning rate  $1 \times 10^{-4}$ .

To help generalize, we decided to increase the dataset to 4000 samples through data augmentation. The hyperparameters of the model were not changed for the second trial, and we obtained training and validation accuracy of 96.6% and 78.5%. However, the training time was longer than an hour, and we realized that accuracies are not improving significantly after epoch 30.




Final Training Accuracy: 0.9475  
Final Validation Accuracy: 0.7883333333333333  
Total time elapsed: 1868.63 seconds  
Test Accuracy: 0.7783333333333333

Figure 8: Training and validation accuracy curves with 4000 samples, 30 epochs, learning rate  $1 \times 10^{-4}$ .

In the third trial, we decided to use 30 epochs, which took 31 minutes for training and validation. The training and validation accuracy was 94.75% and 78.83%, which was not too different from the result we obtained in the second trial. This final model produced a test accuracy of 77.83%.

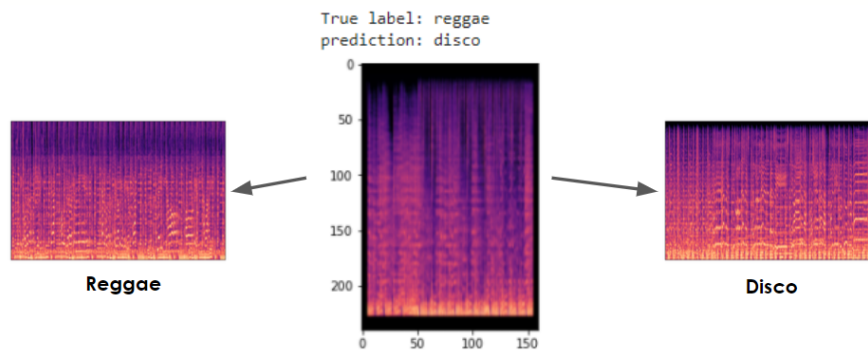
## 8. Qualitative Results

In order to check if its prediction is correct, the model has compared the actual label of the song and predictions, and increased the number of correct predictions when the actual label corresponds to predictions.




*Figure 9: Sample prediction with correct classification.*

This image shows sample prediction with correct classification. It is labelled as classical, and prediction of our model was aligned with the label. To compare it with another mel-spectrogram from classical, they look similar to each other. Both of them are having some black area at the above, followed by a long range of purple and orange, and a short range of orange below. They also showed similar patterns.



*Figure 10: Sample prediction with wrong classification.*

This image shows sample prediction with wrong classifications. It was labelled reggae, but the model predicted it as disco. Comparing it with examples from reggae and disco, all three were showing similar patterns, but reggae does not have black region at the top, while disco does. The image had a black region on top, so it appears like the music clip had some disco traits which could confuse our model.



*Figure 11: Confusion Matrix.*

According to Figure 11, our model has predicted classical and metal clips accurately, while it showed lower accuracy on hip-hop and rock. The most inaccurate prediction was predicting hip-hop as reggae (26%), followed by predicting rock as country (22%), blues as country (11%), and disco as reggae (10%). It confirms that the model could be confused by music clips which show traits of different genres.

## 9. Evaluate Model on New Data




Figure 12. Confusion matrix from new dataset

To test our model on a completely new dataset, we created a new dataset consisting of 3 songs each for each genre, which is 30 songs in total. In the new dataset, we obtained an accuracy of 61.725%. Even though it is slightly lower than test accuracy in [Figure 11], it is still considered better than the baseline test accuracy of 50.67% provided by K-nearest Neighbours algorithm. For this new dataset, hip-hop and rock are the two genres that were predicted most accurately, while our model doesn't perform as well when trying to identify country, metal and reggae genres.

The accuracy from the new dataset is lower than validation or test accuracy, and it can be a consequence of having “bad” data. GTZAN is the dataset that is popular for machine learning projects involving music audio, because it is processed and labeled properly. We used GTZAN datasets for training our models and hyperparameters tuning. However, for the new dataset we obtained our data from other sources on the internet, which has risks of data not being labelled properly or data that are of low quality. If the music clips are not labelled properly, then there is a chance that the accuracy from the dataset is lower. For example, if a song is labeled as “metal” when in reality it is more similar to being classified as “rock”, then this would result in a wrong prediction from our model if our model classifies it as “rock”. Since music can have traits of multiple genres, labelling a song into one genre may not represent the song properly. If we chose songs which show traits of multiple genres, then it is likely that the model would not classify the song's genre properly.

## 10. Discussion

As mentioned in the introduction, we believe that genre-detection is not a trivial task. There are large diversities between genres itself, caused by differences in musicians and creative liberties taken by certain composers. Despite the patterns formed by musical properties within genres, these factors often cause ambiguity in mel-spectrograms and cannot be clearly parsed by the machine. For example, in [Figure 10], the mel spectrogram was labelled as reggae, but our model classified it as disco, since it had similar characteristics with disco's mel-spectrogram. This can justify our low testing accuracy, and may also justify the existence of a performance plateau of genre-detection machines in general.

One of the more important things we learned was the importance of having large datasets in complex problems, as it can greatly improve a model's accuracy. We were able to increase validation accuracy by 24.5% after creating a larger dataset using data augmentation. Since data augmentation was done by simply manipulating existing data to create new data, we didn't expect large improvements in performance. The new data, however, introduced enough diversity to allow the model to generalize efficiently.

## 11. Ethical Consideration

Music is a form of artistic and creative expression. There isn't a definite description of how a certain genre of music has to sound like. A song can be very dynamic, composed of elements from many genres of music. By trying to give a definite answer and classifying a song to be only one genre of music, we are restricting its artistic expression. This puts a preconceived notion on a song that may go against the ideas of its creator. For example, our model may classify a song to be Rock music, but its author may claim otherwise.

Most music streaming platforms have playlists of different music genres. If a song is classified as only one genre and is only placed in one of the many genre-specific playlists, it will lose many potential listeners who don't listen to other playlists because of its labeled genre. As a result, an artist may lose profits on music royalties or the reach of a bigger audience. Furthermore, there is a potential for our project to replace jobs in the music industry that involves music classification, like playlist curators on music streaming platforms.

## 12. Project Difficulty

Throughout the report, we have repeatedly noted how the difficulty in identifying genres stems from the ambiguous similarities between genres. For example, the two genres rock and metal are often misclassified by even human ears, a decision that usually comes down to personal opinion.

We attempted to thwart this inherent and unavoidable property by including an LSTM module to learn from musical arrangement, as humans don't actively consider arrangement when determining a genre. With a fairly deep network and the LSTM in place, we only managed to acquire a 77.83% accuracy.

At the start of hyperparameter tuning, while attempting to deepen the network, we noticed that we were not only overfitting to the training dataset but severely slowing down the training speed. To mitigate this, we worked with a smaller network and implemented dropout layers to help generalize the model, and batch normalization layers for training efficiency.

We only managed to achieve a validation accuracy of 50% with this architecture, and decided that the only way to improve generalization was to generate more data. With a limited audio file data set, we used data augmentation techniques of noise injection, pitch shifting and time stretching to increase the amount of data that we have to train our network. As a result, we were able to drastically increase the validation performance from 50% to around 75%.



## 13. References

- [1] T. Greer and S. Narayanan, “Using shared vector representations of words and chords in music for genre classification,” SMM19, Workshop on Speech, Music and Mind 2019, 2019.
- [2] D. A. Huang, S. A. Arianna, and P. J. Eli, “Music Genre Classification,” 2018.
- [3] Index of /sound/genres. [Online]. Available: <http://opihi.cs.uvic.ca/sound/genres/>. [Accessed: 12-Feb-2021].
- [4] “1.4. Support Vector Machines¶,” scikit. [Online]. Available: <https://scikit-learn.org/stable/modules/svm.html>. [Accessed: 12-Feb-2021].
- [5] “sklearn.neighbors.NearestNeighbors¶,” scikit. [Online]. Available: <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.NearestNeighbors.html#sklearn.neighbors.NearestNeighbors>. [Accessed: 12-Feb-2021].
- [6] “Lecture 2: k-nearest neighbors / Curse of Dimensionality,” Cornell Computer Science - CS4780. [Online]. Available: [https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote02\\_kNN.html](https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote02_kNN.html). [Accessed: 12-Feb-2021].
- [7] R. Berwick, “An Idiot’s guide to Support vector machines (SVMs),” in 6.034 Artificial Intelligence - MIT.